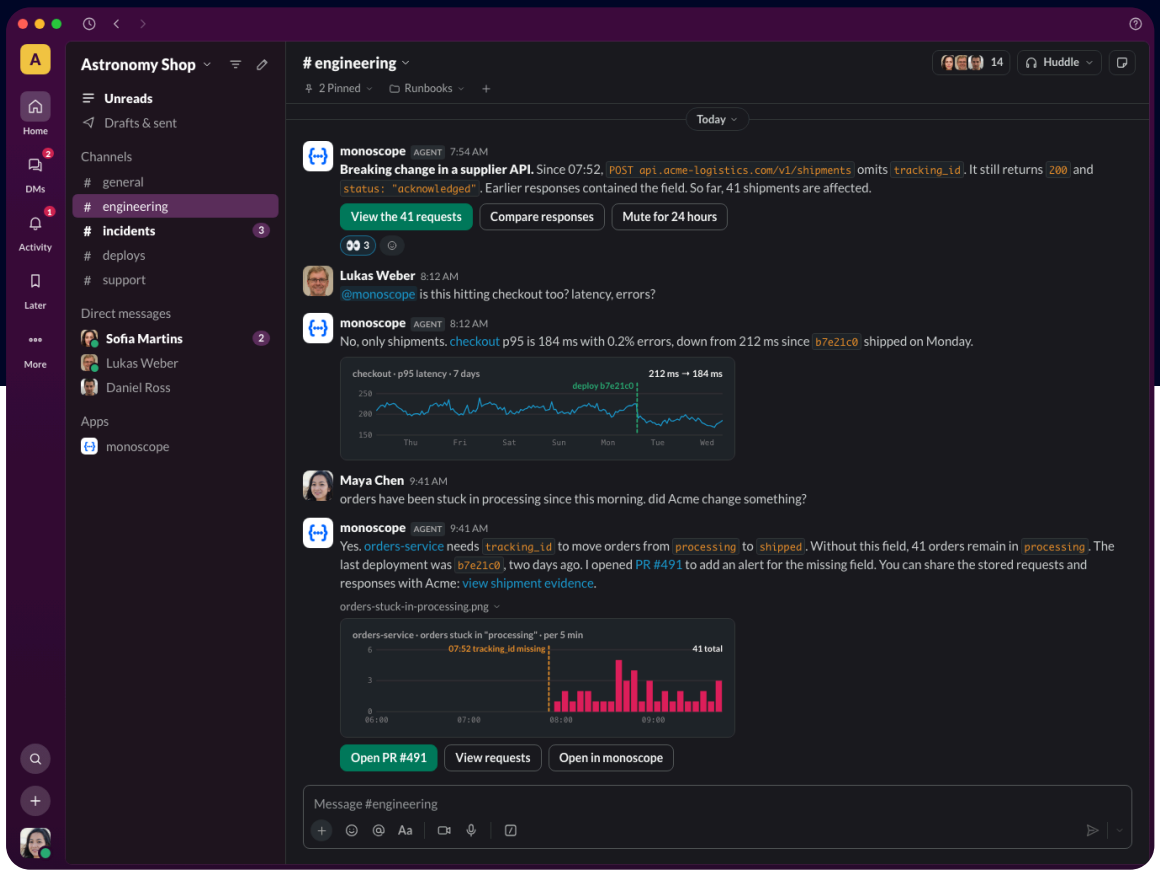
Know what happened. Prove it. Fix it.

When something breaks, watch a session replay of what the user saw and did. Some issues are easier to see than explain. The exact request, response and trace sit alongside the replay. monoscope's AI investigates, reports in Slack and opens a fix for review.



07:52

A logistics supplier quietly stops sending `tracking_id`, but its API keeps returning **200 OK**, so nothing looks wrong.

07:54

monoscope spots the breaking change two minutes later and posts it in Slack, with all 41 affected requests attached.

09:41

When an engineer asks what's going on, monoscope answers from the stored requests and opens PR #491.

andela Platnova PayFonte Grovepay SAMEDAY CUSTOM thepoor woodcore Blockradar

A 200 OK that stranded 41 orders

Your trace will happily show a fast, successful call. But if a supplier drops a field and you never kept the response body, there's nothing to explain why your customer's order is stuck.

What a typical trace keeps

`POST /v1/shipments` `200` `182 ms`

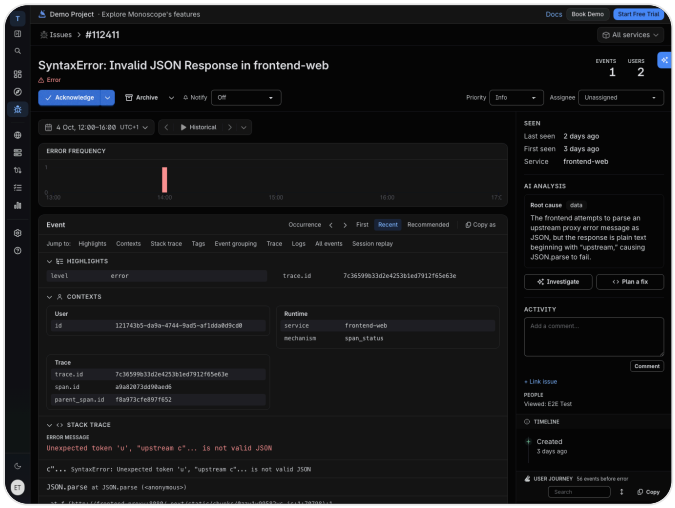
- ✕ No request body
- ✕ No response body
- ✕ Possibly not this request at all, if it was sampled out
- ✕ Nothing you can send the supplier as proof

What monoscope keeps for every request

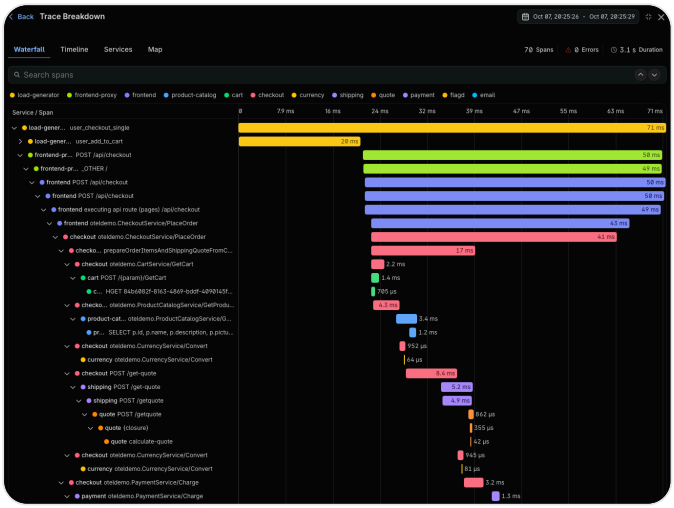
Response from api.acme-logistics.com at 07:52:14.208

```
HTTP/1.1 200 OK
{
  "shipment_id": "shp_7Hk2m9",
  "status": "acknowledged",
  "tracking_id": "AX-99812-NG",   missing since 07:52
  "eta": "2026-10-09"
}
```

monoscope notices the missing field, flags it as a breaking change and links it to the trace, the logs and the customer's session.



Related errors are grouped into one issue, with the evidence and an AI analysis already attached.



From the issue, you're one click from the full trace and every service the request touched.

Never sampled

monoscope keeps every request, response, log and span you send, so you never have to guess which events were worth keeping.

Evidence you can share

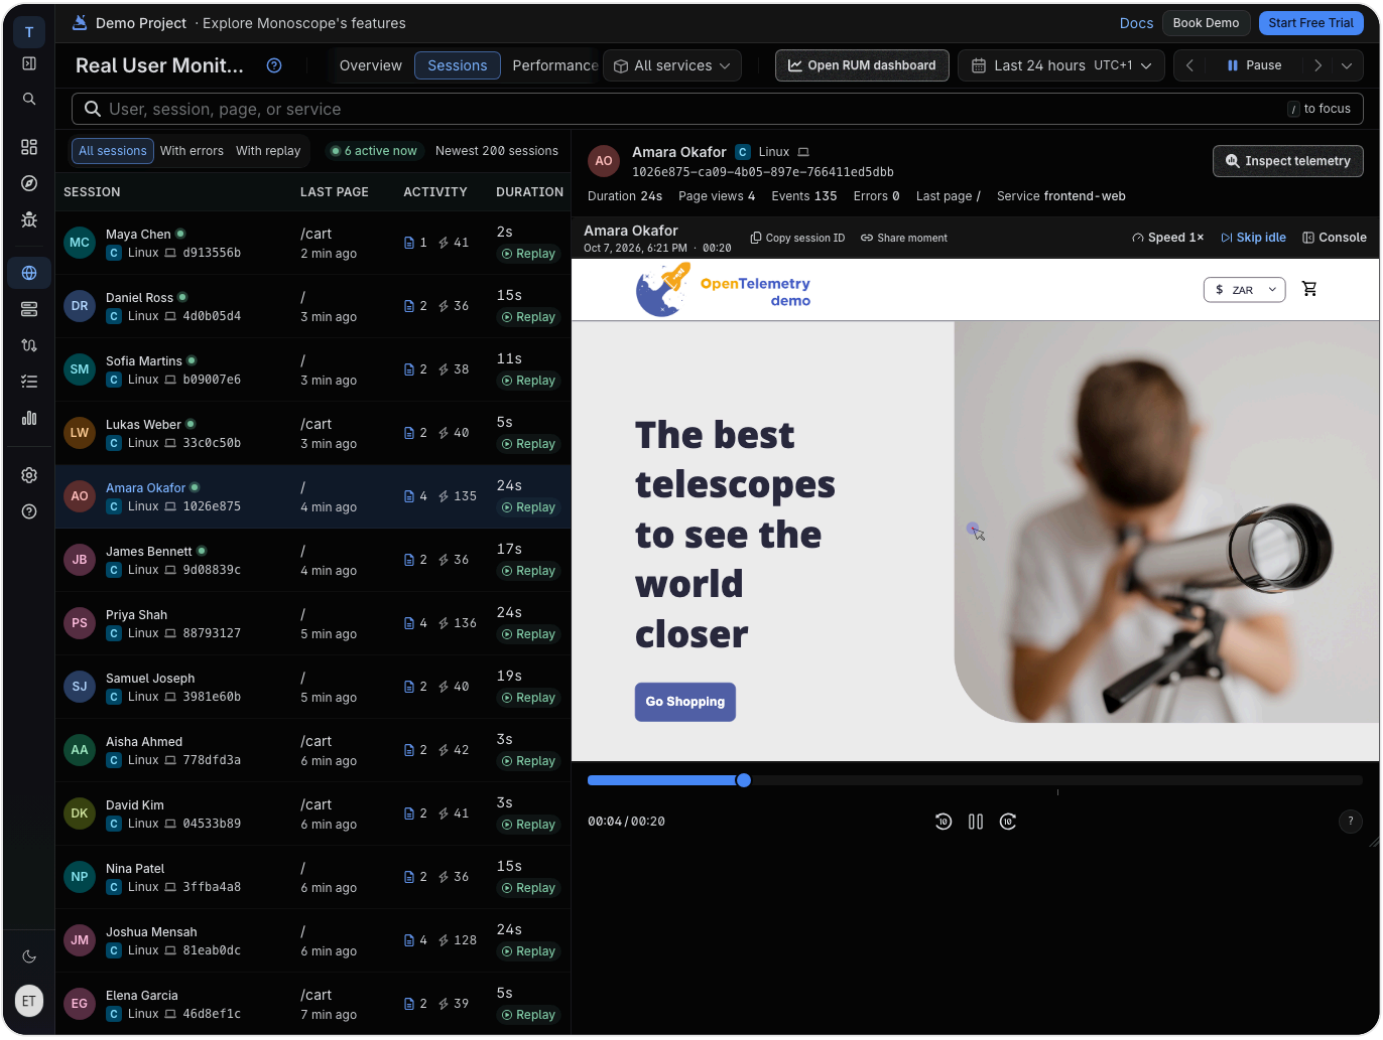
You can see exactly what you sent, what the supplier sent back and when, down to the millisecond, and share it all with a link.

Your bucket is the database

On Cloud + your own S3, your data lives in your own bucket as open Delta Lake tables, and you can keep it for as long as you like.

Some issues are easier to see than explain.

Watch a session replay of what the user saw and did when an issue occurred. Then inspect the request, response and trace.



Watch the user's screen. Pause the replay, inspect their actions and open the session's telemetry. Recorded demo session from the running product.

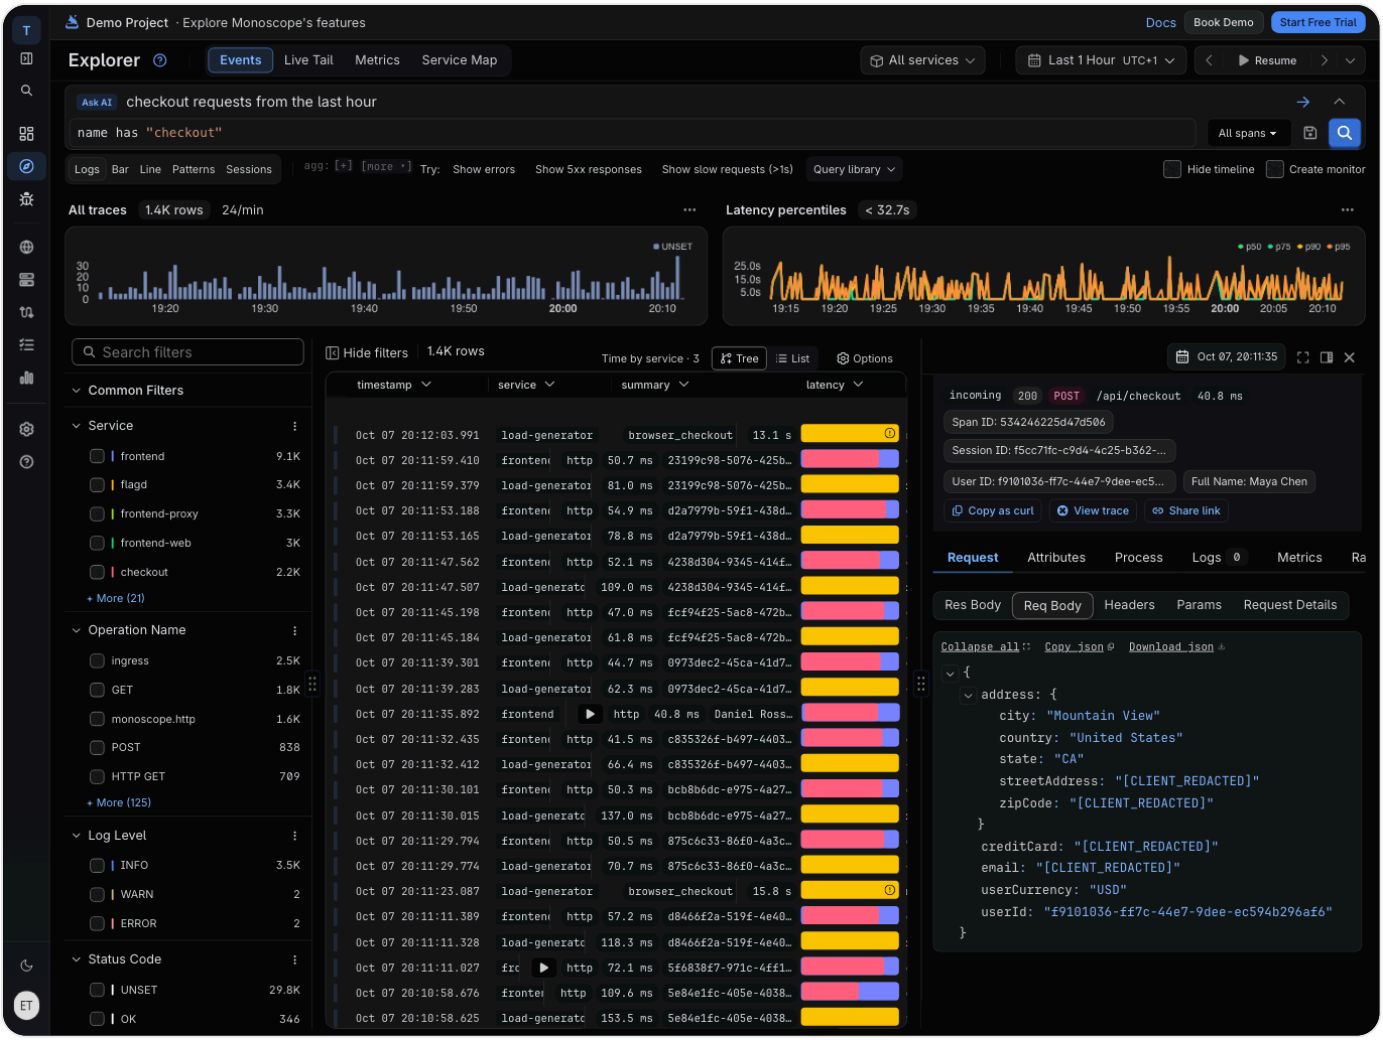
See the sequence

Watch the user's actions and the screens they saw. Pause or replay the moment you need to inspect.

Connect it to the evidence

Open the session's telemetry to inspect the requests and traces behind the user's experience.

What your app sent, and what came back.



Ask in plain language, then read the request. monoscope writes the query and shows it. The body your app sent is there, with sensitive fields redacted.

Read both bodies

A status code tells only part of the story. Open the request and the response, stored in full.

Find the missing field

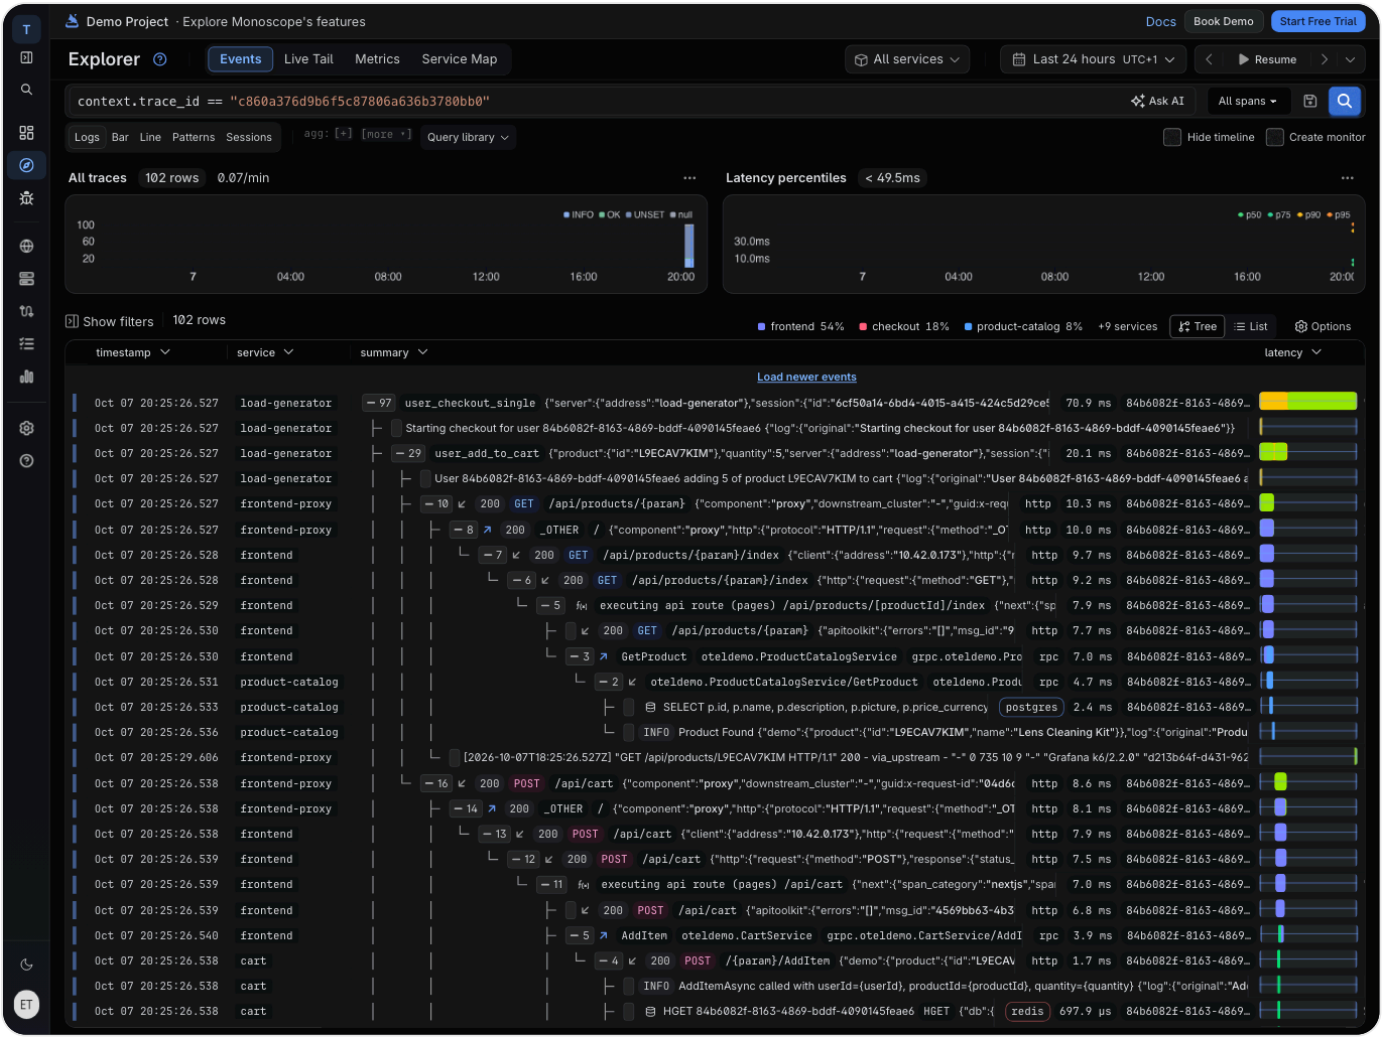
Compare the response with what your app expected, and see why a 200 still caused an issue.

Share the evidence

Send the stored request and response to your team or the supplier with a link.

One checkout, everything it touched.

Open a trace in the explorer and every span and log nests under the request that started it: HTTP calls, gRPC, SQL, Redis and the logs in between.



102 events from one checkout, across 12 services, in the order they happened.

Spans and logs together

Logs sit under the span that wrote them, so you read them in context, not in a separate tool.

Every query and cache call

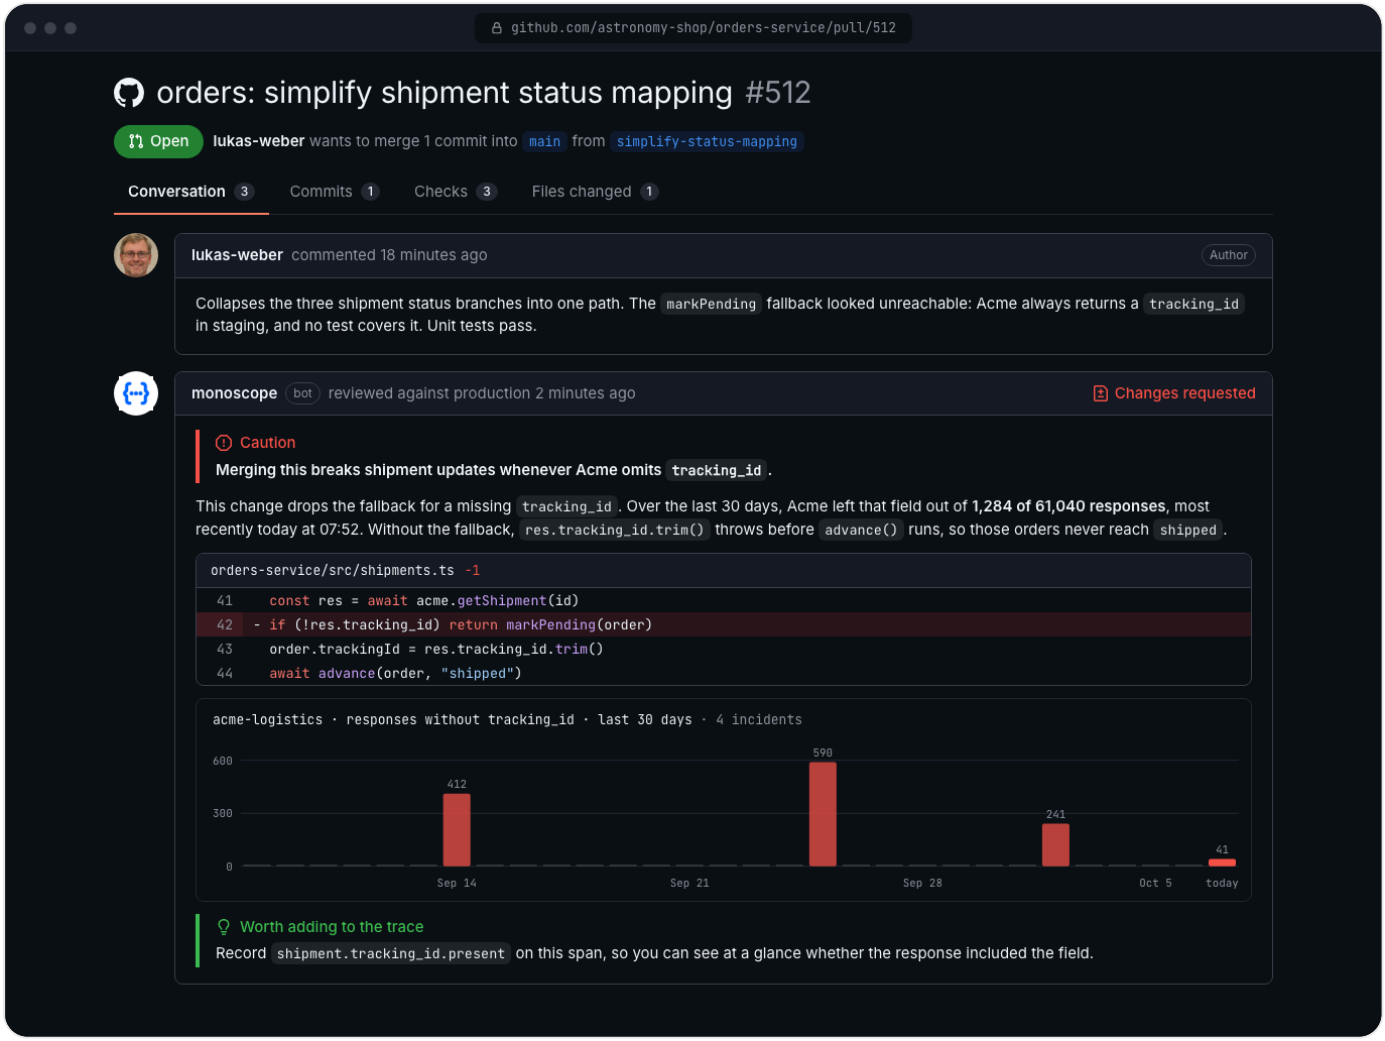
See the SQL a service ran and the Redis keys it read, with timings for each.

Nothing sampled out

monoscope keeps every event, so the trace you need is complete when you open it.

Find the next incident before you merge.

monoscope reviews every pull request, from engineers or coding agents, against your live production data, and shows the evidence when a change looks risky.



Answers in Slack

Ask why checkout is slower, and get an answer from your logs, traces and deploys.

Remembers past issues

Recurring errors keep their history, so each investigation builds on the last.

Finds monitoring gaps

monoscope points out endpoints nobody's watching and ranks alerts by impact.

Works with your agents

Plug it into your own agents through a hosted MCP server with about 50 tools, an agent-mode CLI or Claude Code skills.

Shows its evidence

Every finding links back to the requests, traces and deployments behind it, so you can check the reasoning before you merge.

Reads the actual request

The AI works from the full request and response, along with past events and your deployment timeline, not just an error message.

What our customers have seen

20x	70%	9x	5x	3x	500ms
faster mean time to resolution	less manual troubleshooting	fewer customer support calls	faster incident debugging	faster system recovery	less response time on improved endpoints
Blockradar	SAMEDAY CUSTOM	Platnova	CORONA MANAGEMENT SYSTEMS	partna	PayFonte

"We had a major incident, and our tech support could see via [monoscope] which third-party integration partner was responsible, and could take action without needing the engineering team's help."

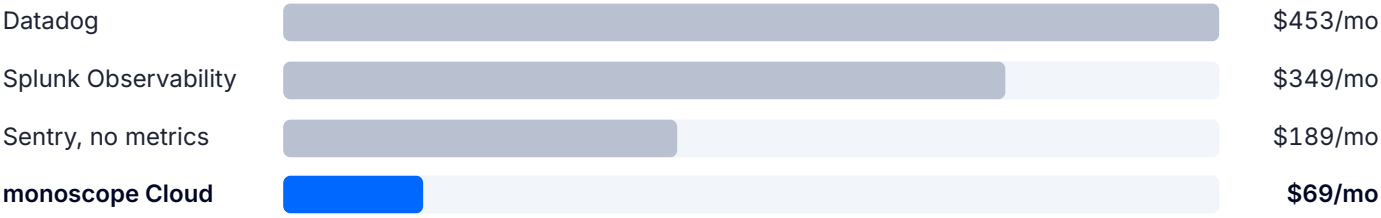
Joshua Chinemezum, CEO, Platnova

"The best observability tool we use today at Woodcore, monoscope notifies us about any slight change that happens on the system."

Samuel Joseph, CEO, Woodcore

Pay per event, not per host

You pay the same rate for requests, logs and spans, with separate allowances for metrics and session replay. There are no per-host fees, so adding servers won't change your bill.



Based on an example workload of 50M events a month, 5 hosts, 2,000 metric series and 5,000 replay sessions, using September 2026 list prices and 30-day retention. You'll find the full assumptions on page 8.

1. **Send your telemetry**
   Add our SDK, or point your existing OpenTelemetry exporter at monoscope. We support 780+ integrations.
2. **Run it side by side**
   Keep your current tool running for a week or two and compare which errors and API changes each one catches.
3. **Move your services**
   Once you're happy, send the rest of your services over and retire the tools you don't need anymore.

See it on your own traffic

You can start free with up to 10,000 events a day, no credit card needed. Or book a call, and one of our engineers will help you get set up.

Start for free
app.monoscope.tech

Try the playground
no signup

Book a call
30 minutes

Where your data lives, and who's in control of it

	Cloud	Cloud + your own S3	Self-hosted
Data stored in	monoscope-managed storage	Your S3-compatible bucket, as open Delta Lake and Parquet tables	Your servers
Retention	30 days	As long as you like	Your choice
Price	Free up to 10,000 events/day. Paid Cloud: $29/month for 20M events, then $1 per extra million.	From $199/month for 100M events. Additional events cost $1 per million.	Community Edition is free (AGPL-3.0). Enterprise starts at $500/month.

Security and privacy

GDPR We're happy to sign a data processing agreement (DPA) whenever you need one.

EU data region If you'd prefer, we can store your data in the EU.

SSO Single sign-on and advanced authentication come with Enterprise.

Redaction The SDK strips out the headers and body fields you choose, such as `$.password`, before any data leaves your app.

Data ownership On Cloud + your own S3, your bucket is the database, so you can query it directly with DuckDB, Spark or SQL.

Open source The whole platform is open source on GitHub under AGPL-3.0, so your team can read the code that handles your data.

Startup program

If you're an eligible startup, we'll cover up to $10,000 of monoscope in your first year.

Volume pricing

Sending more than about a billion events a month? Email hello@monoscope.tech and we'll put together custom pricing.

How the cost estimate on page 7 was calculated

We based these estimates on September 2026 public list prices, with annual billing where it's offered and 30-day retention for every vendor.

Workload: 50M events/month, 5 hosts, 2,000 metric series and 5,000 replay sessions. Each event averages 2 kB. Each metric series sends one data point per minute.

Datadog: Infrastructure costs $15/host and APM costs $31/host. Log ingestion costs $0.10/GB. Indexed logs cost $2.50/million. Custom metrics cost $5/100 series beyond 100/host. Per 1,000 sessions, RUM costs $0.15 and Session Replay costs $2.50.

Splunk: Observability Cloud App & Infrastructure costs $60/host. Splunk Cloud log ingestion costs about $12.50/month per GB of daily ingestion. RUM costs $14/10,000 sessions. This estimate excludes custom metrics.

Sentry: Team costs $26/month. Spans cost $1.60/million beyond 5M. Errors cost $0.29/1,000, with an assumed error rate of 0.5%. Replays cost $3.75/1,000. This estimate excludes metrics.

monoscope Cloud: $29/month includes 20M events, 20M metric data points and 2,000 replay sessions. Additional usage costs $1/million events, $1/10M metric data points and $1/1,000 replay sessions.

Want to try your own numbers? Use the pricing calculator on our site.